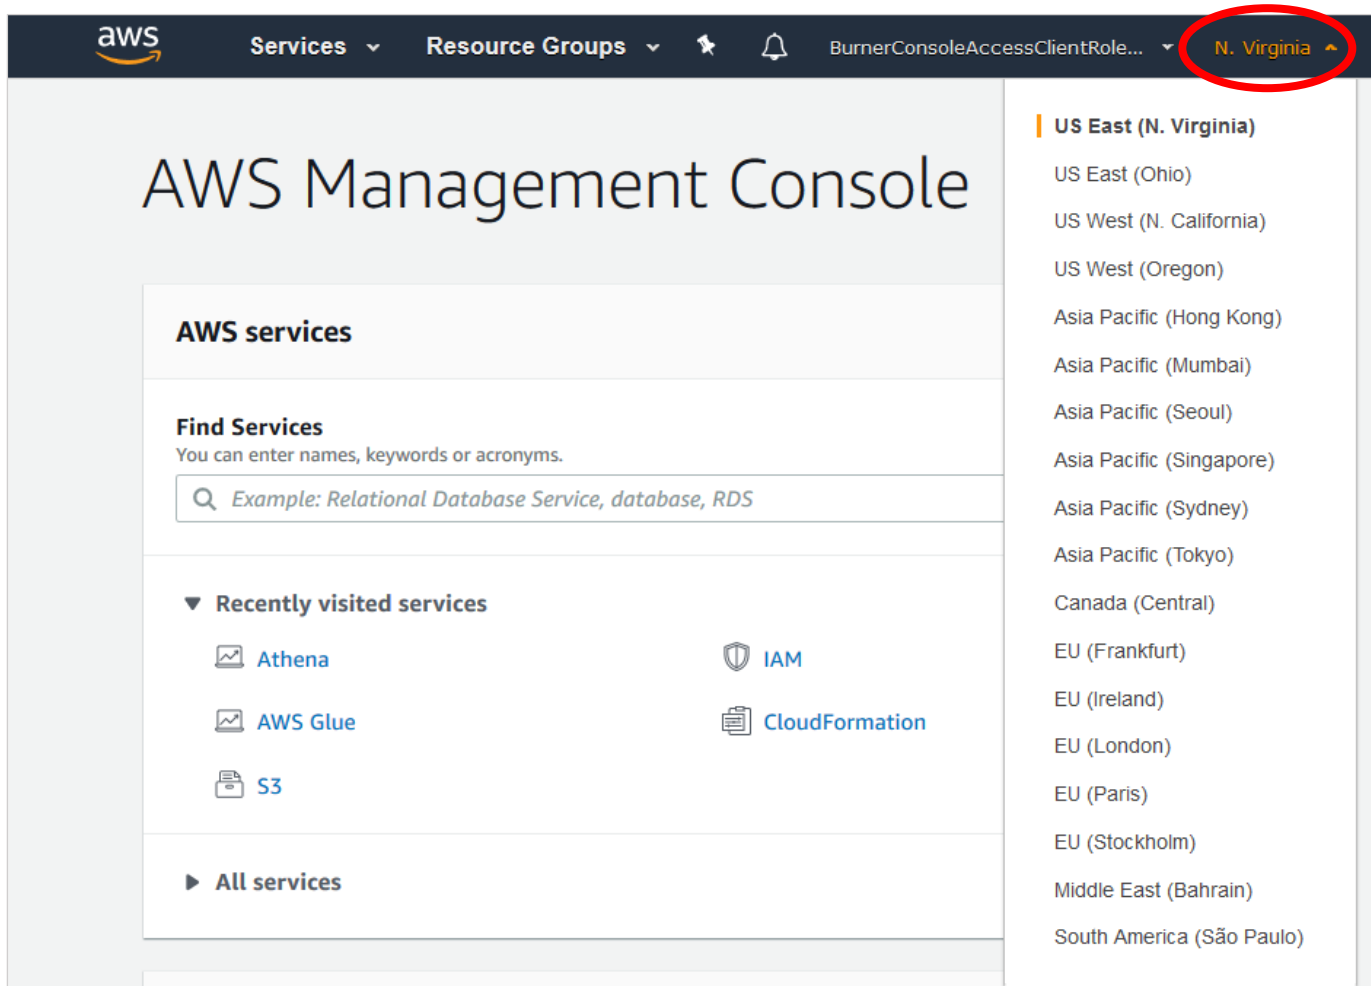


Section 4

Create an S3 bucket and subdirectories as a webserver

In this section, you will create another S3 bucket, this one to host static website content. You will then upload an HTML file to let your users take advantage of the API you just created via the web. You won't have to manage any webserver—your HTML file will be served by S3 bucket configured to act as a web server.

- 1) Log on to the AWS console and change your region to N. Virginia.



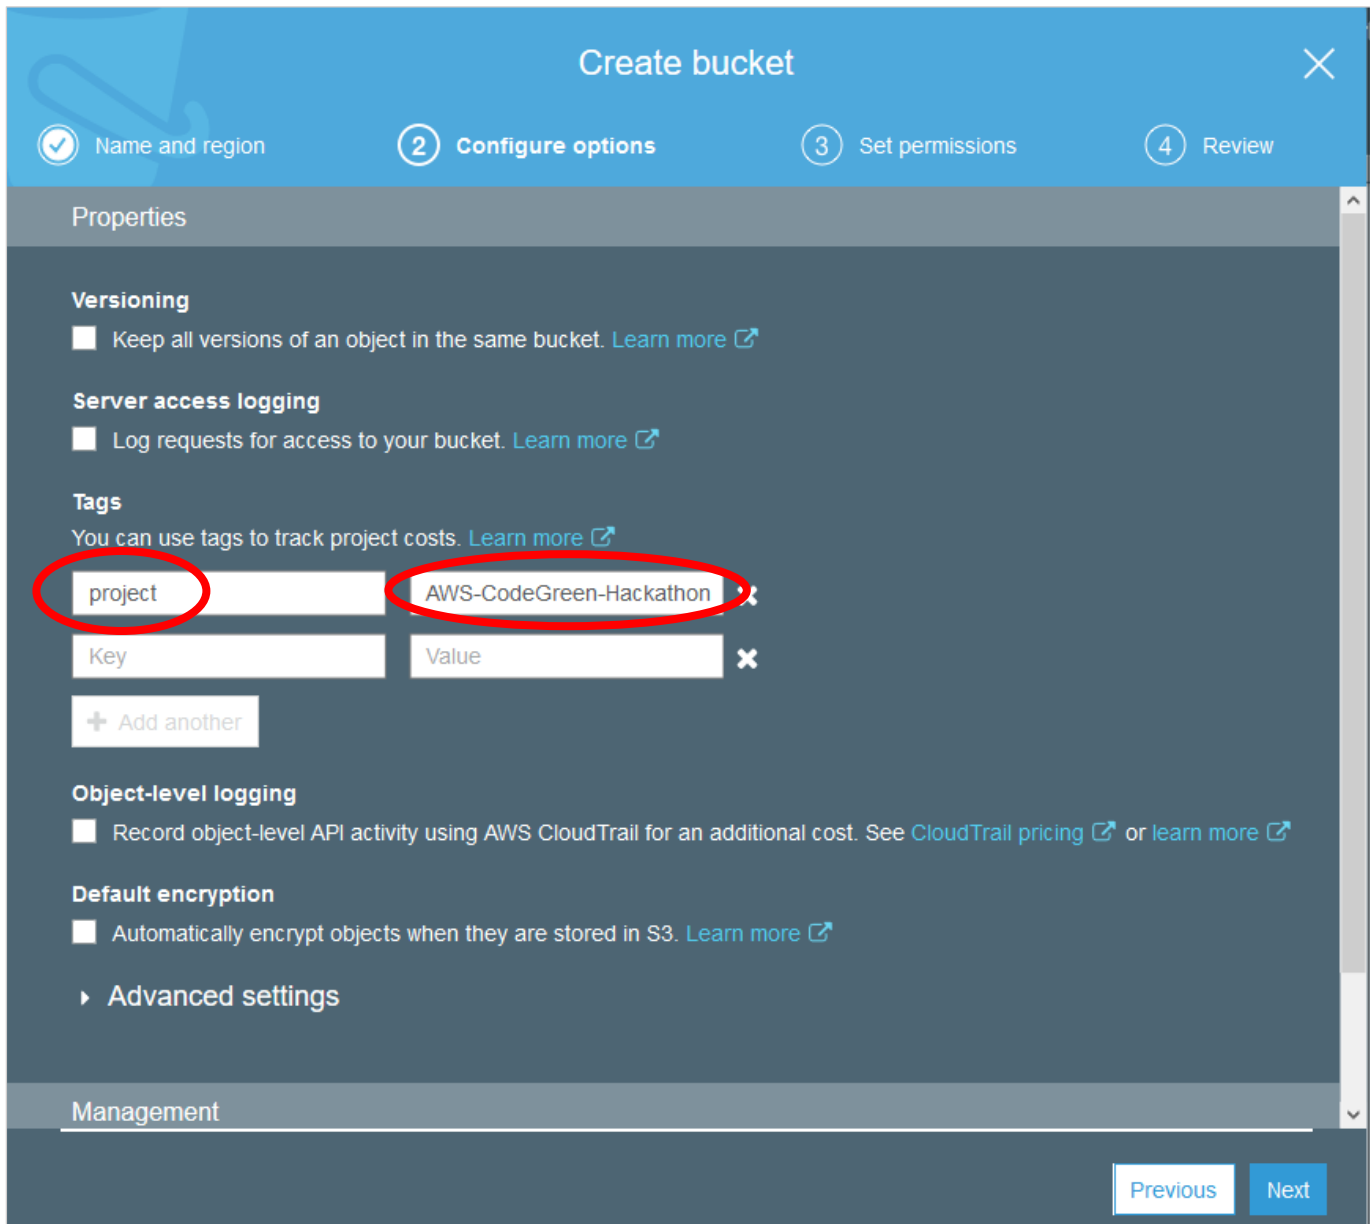
- 2) In the "Find Services" field, search for S3, navigate to the S3 dashboard, then click "Create bucket". All S3 buckets must have a globally unique name and must comply with DNS naming conventions; generally use lower-case letters and no underscores ([more information](#)). We recommend using your initials-web or some other name for uniqueness. Select the region "US East (N. Virginia)" and a unique name for the



bucket name and click "Next".

The screenshot shows the 'Create bucket' wizard in the AWS Management Console. The first step, 'Name and region', is active. The 'Bucket name' field contains 'rcr-web' and the 'Region' dropdown is set to 'US East (N. Virginia)'. Both fields are circled in red. Below these fields is a section for 'Copy settings from an existing bucket' with a dropdown menu showing 'You have no buckets0 Buckets'. At the bottom, there are three buttons: 'Create', 'Cancel', and 'Next'. The 'Next' button is highlighted in blue, indicating it is the next step in the process.

- 3) Add a tag with "project" as the key and "AWS-CodeGreen-Hackathon" as the value, and click "Next".



Create bucket

1 Name and region 2 **Configure options** 3 Set permissions 4 Review

Properties

Versioning
☐ Keep all versions of an object in the same bucket. [Learn more](#)

Server access logging
☐ Log requests for access to your bucket. [Learn more](#)

Tags
You can use tags to track project costs. [Learn more](#)

Key	Value
project	AWS-CodeGreen-Hackathon

[+ Add another](#)

Object-level logging
☐ Record object-level API activity using AWS CloudTrail for an additional cost. See [CloudTrail pricing](#) or [learn more](#)

Default encryption
☐ Automatically encrypt objects when they are stored in S3. [Learn more](#)

[Advanced settings](#)

Management

[Previous](#) [Next](#)

- 4) Since this bucket will be acting as our webserver, public access to this bucket is required. Uncheck "Block all public access" and click "Next".

Create bucket

✓ Name and region

✓ Configure options

3 Set permissions

4 Review

Block public access (bucket settings)

Public access is granted to buckets and objects through access control lists (ACLs), bucket policies, or both. In order to ensure that public access to all your S3 buckets and objects is blocked, turn on Block *all* public access. These settings apply only to this bucket. AWS recommends that you turn on Block *all* public access, but before applying any of these settings, ensure that your applications will work correctly without public access. If you require some level of public access to your buckets or objects within, you can customize the individual settings below to suit your specific storage use cases. [Learn more](#)

☒ **Block *all* public access**

Turning this setting on is the same as turning on all four settings below. Each of the following settings are independent of one another.

☐ **Block public access to buckets and objects granted through *new* access control lists (ACLs)**

S3 will block public access permissions applied to newly added buckets or objects, and prevent the creation of new public access ACLs for existing buckets and objects. This setting doesn't change any existing permissions that allow public access to S3 resources using ACLs.

☐ **Block public access to buckets and objects granted through *any* access control lists (ACLs)**

S3 will ignore all ACLs that grant public access to buckets and objects.

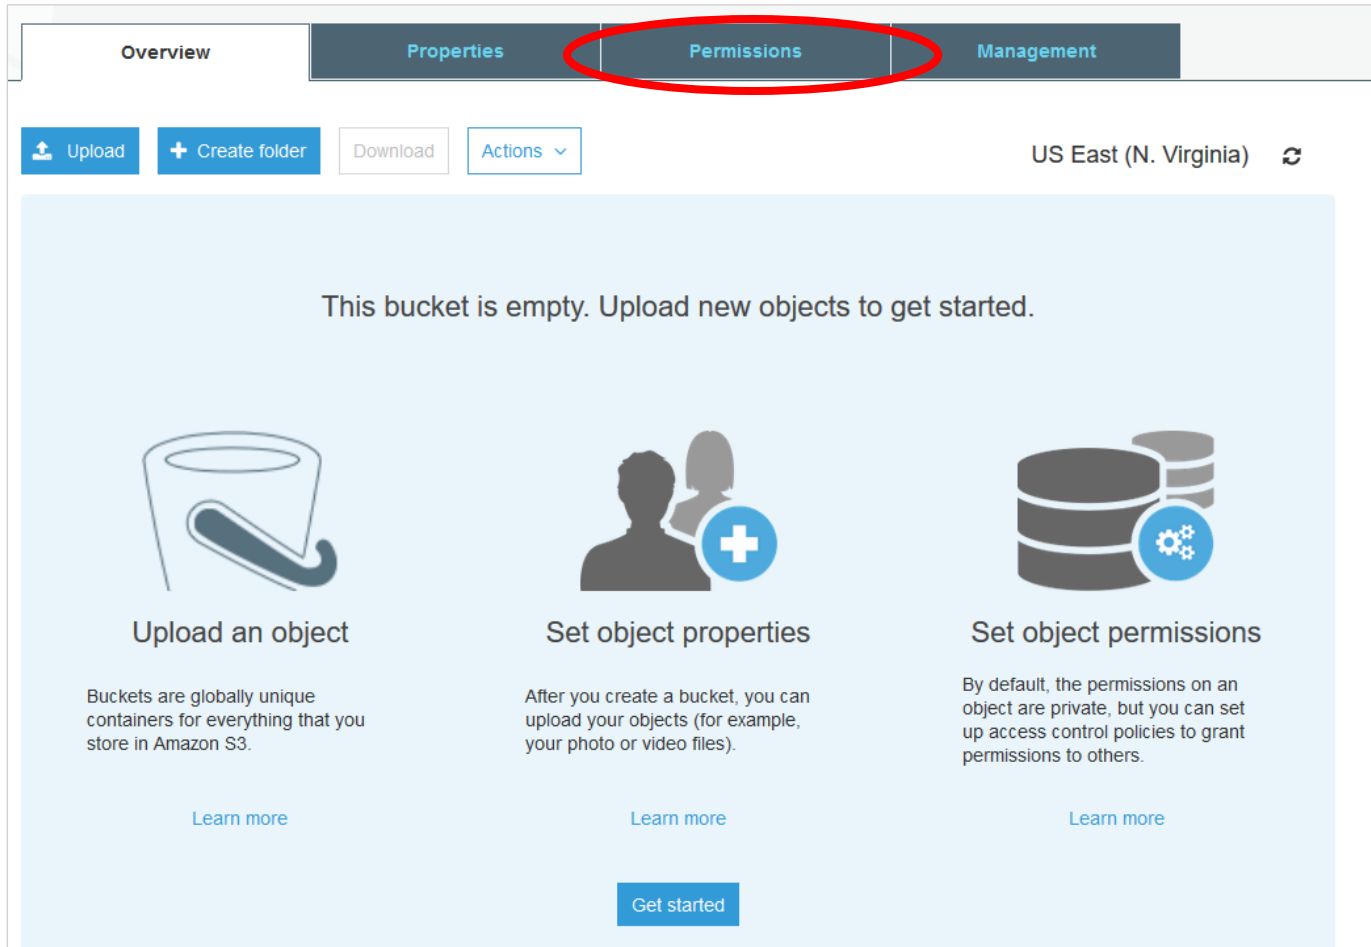
☐ **Block public access to buckets and objects granted through *new* public bucket policies**

S3 will block new bucket policies that grant public access to buckets and objects. This setting doesn't change any existing policies that allow public access to S3 resources.

Previous

Next

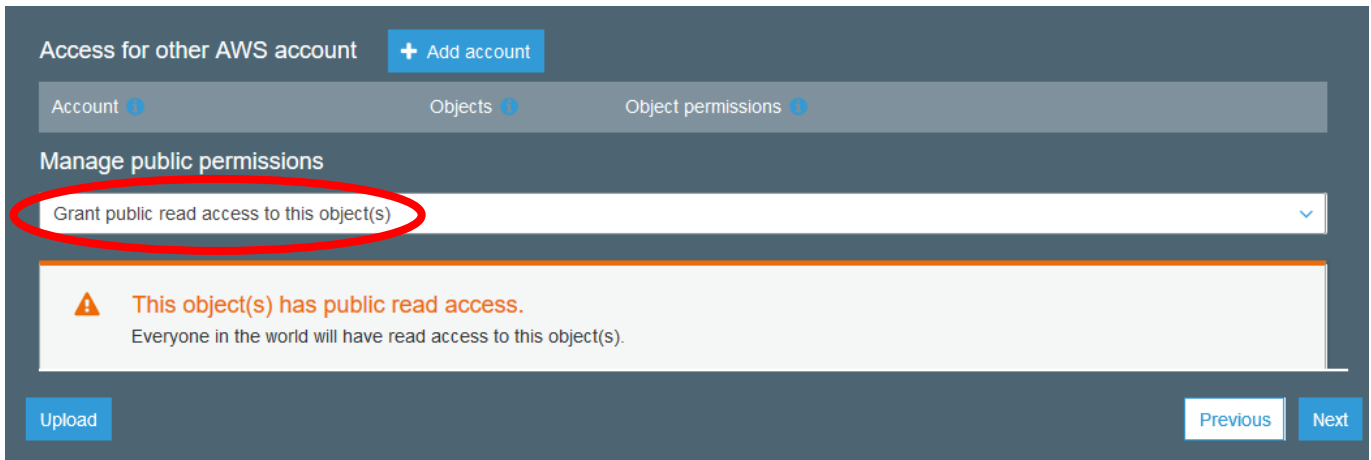
- 5) Click on "Create bucket".
- 6) Now you will see your bucket on the Amazon S3 dashboard, under Buckets. Notice that under "Access", your bucket is listed as "Bucket and objects not public". To make the objects in your bucket publicly readable, you must write a bucket policy that grants everyone the `s3:GetObject` permission. The sample bucket policy shown in the following steps grants everyone access to the objects in the specified folder.
- 7) Click on the bucket you just created, then click on the "Permissions" tab.



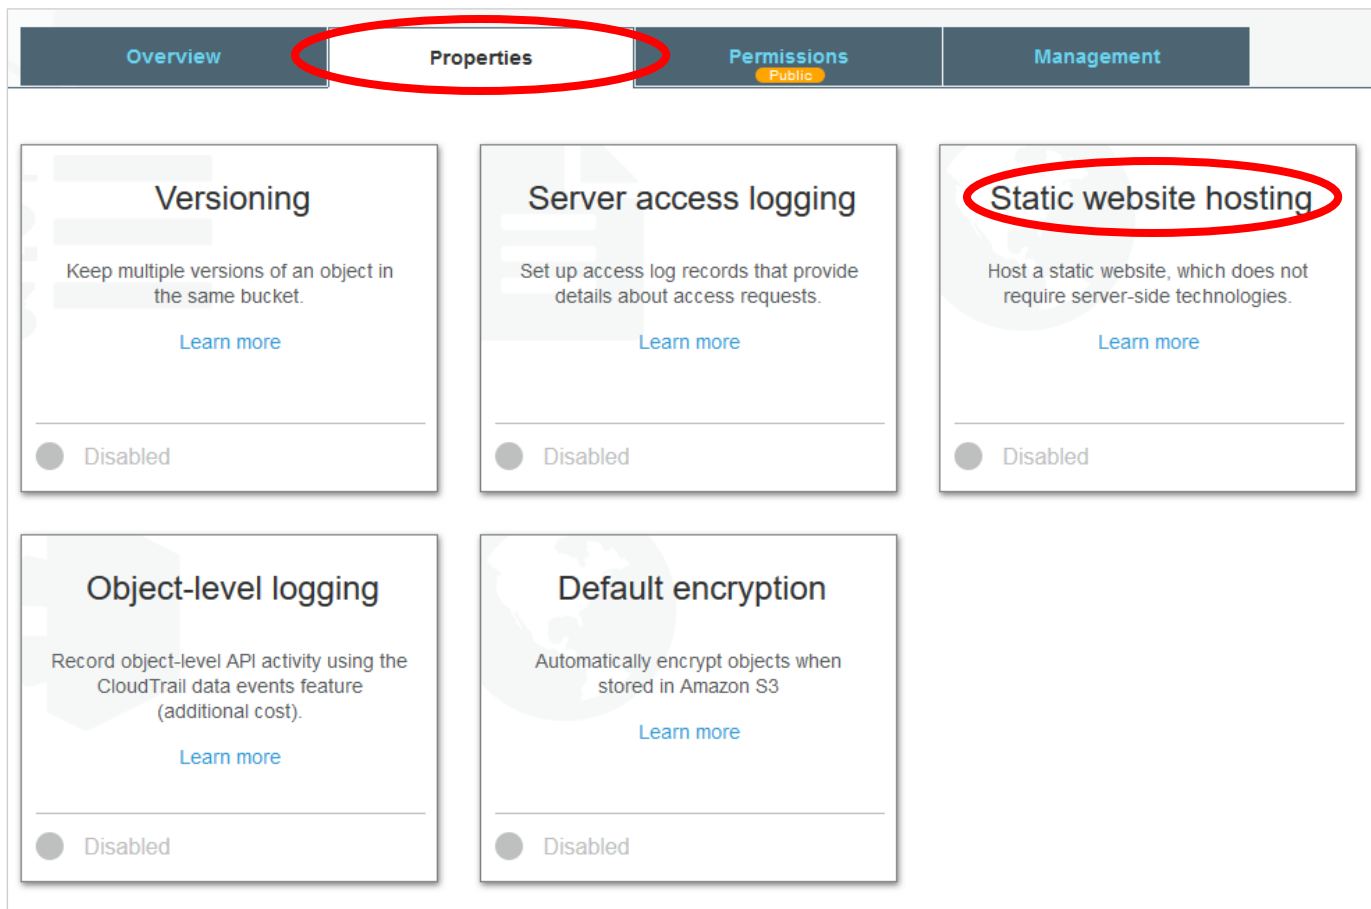
- 8) Click on the “Bucket Policy” button. Open the [S3-policy-web.json file](#) and replace **YOUR-BUCKET-NAME-HERE** with the name of the bucket you created earlier in the exercise, then click “Save”.

```
{
  "Version": "2012-10-17",
  "Statement": [{
    "Sid": "PublicReadGetObject",
    "Effect": "Allow",
    "Principal": "*",
    "Action": ["s3:GetObject"],
    "Resource": ["arn:aws:s3:::YOUR-BUCKET-NAME-HERE/*"]
  }]
}
```


- 9) You will be given a warning about making your S3 buckets public. Keep in mind, anyone on the Internet will be able to READ data in this bucket.
- 10) Click the "Overview" tab, then click on the "Upload" button. Upload the index.html file from the "code" directory, which you can find [here](#). Click "Next".
- 11) In the Public permissions, click on the drop-down box and select "Grant public read access to this object(s)" and click "Next".



- 12) Accept the defaults for standard Storage class and click "Next".
- 13) Click "Upload".
- 14) Now click on the "Properties" tab, then click "Static website hosting".



- 15) Copy the endpoint URL and save it to your text editor, as this will serve as the website URL you are creating. Select the "Use this bucket to host a website" radio button, type index.html for the Index Document and click "Save". We are not configuring an error document, but you could define it here if you decide to create one on your own.
- 16) Open a new tab in your web browser and go to the endpoint you copied in the previous instruction.
- 17) Click "Submit", wait 15-20 seconds, and this is the ideal location to run your event!
- 18) The response you just got was from a previously created endpoint—that's what index.html is using out of the box. Now we're going to update the app to use the endpoint you just created. First, click "use your own endpoint":



- 19) Paste in the endpoint you copied above, and re-run your query. Your app is now using the endpoint (and API Gateway, Lambda, and Athena table) you just created!

here.'" data-bbox="139 173 784 311"/>

In this section, you created an S3 bucket to serve the static website which utilizes the application you built over the course of this workshop.

You've now completed the 4th and final section of the workshop—congratulations! Now that you're done, explore some optional [next steps and useful links](#) to see how you can build on what you've created here.